

# Finite Element Mesh Partitioning Using Perl+OpenMP



The Perl & Raku Conference  
July 11-14, 2023  
Toronto, CA Eh!



# A Perl Boot Camp



# About Perl Camp

# Goals for this Talk Are Not

- ◆ See pretty code
- ◆ See performance that scales as threads increase
- ◆ Present that OpenMP is a silver bullet\*

\* that's my  
term!



# Goals for this Talk Are

- ◆ Increase awareness of OpenMP in GCC
- ◆ Convince the Perl Community of the enormous opportunity embracing GCC's OpenMP represents
- ◆ Begin to identify key obstacles to more widespread awareness and adoption
- ◆ Present OpenMP as an available bullet\* in Perl

\* make my  
day!



# *Let's Define the Target User*

## **C/OpenMP User, Perl dabbler**


- ◆ Wants rapid prototyping of Perl
- ◆ Wants access to easy work-sharing offered by OpenMP
- ◆ Doesn't want to learn Perl API
- ◆ Probably has existing C/OpenMP code they wish to embed in a Perl driver
- ◆ Domain scientist (physical sciences, DMKD/ML)
- ◆ Not ready for PDL

## **Perl User, C dabbler**

- ◆ Has already decide Perl is part of the solution
- ◆ Has tried other “thread” approaches in Perl, has found nothing in the *Goldilocks zone*
- ◆ Has been exposed to HPC or has supported HPC users
- ◆ Wants to process large datasets with lots of computational *locality*
- ◆ Has a machine with 8+ cores

# DBSCAN

A data partitioning algorithm using a density function



<https://dev.to/wmisingo/machine-learning-text-clustering-with-dbscan-399n>

# ADCIRC

Predicts hurricane storm Surge using an  
“unstructured” mesh (triangular elements)

<https://adcirc.org>




Figure E1.1 Geometry of the quarter annular test problem showing the location of the elevation and velocity stations (stations 1, 2, and 3)





Figure E1.2 Finite element discretization of the quarter annular test problem showing node and element numbers





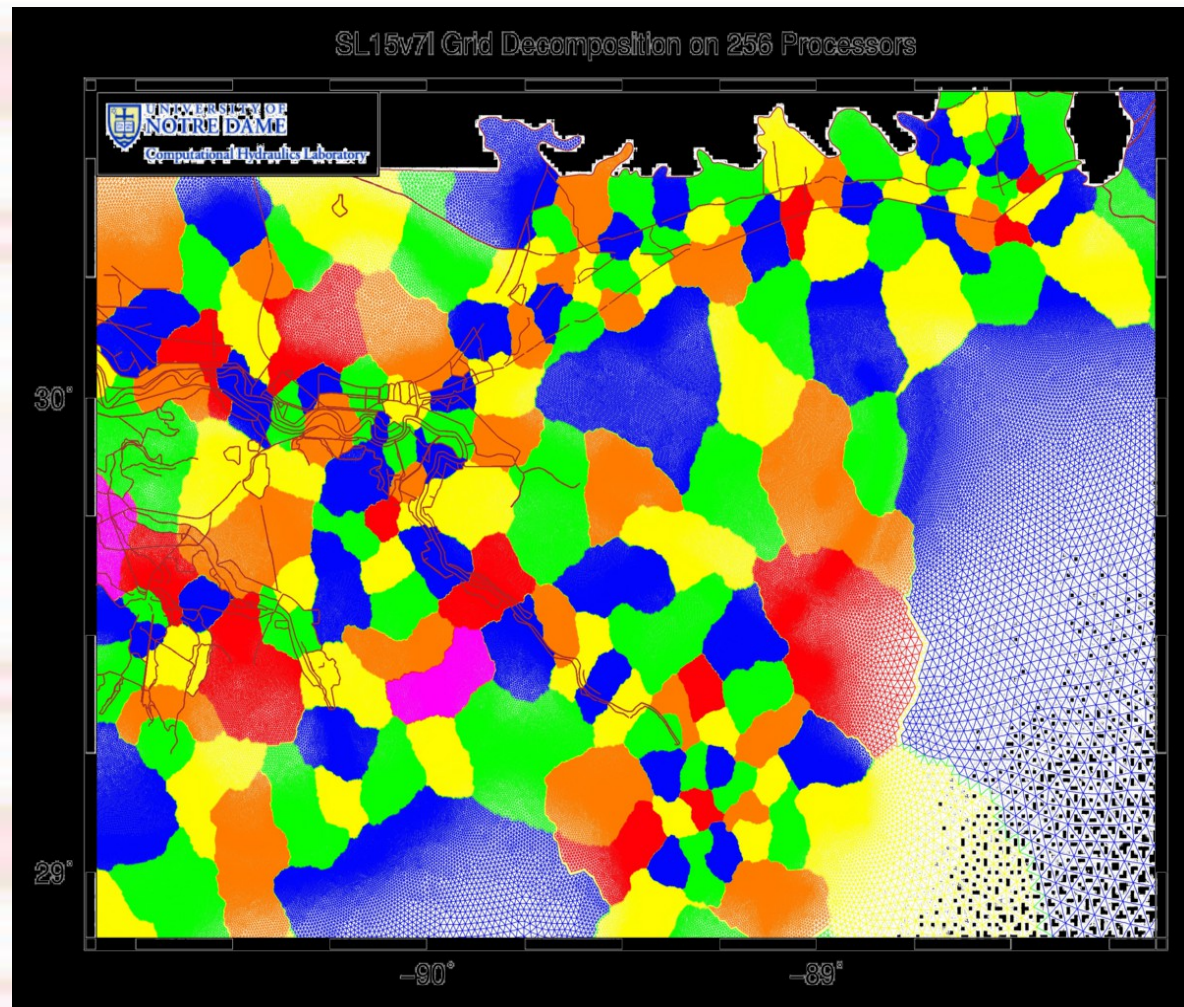
# Parallel ADCIRC

Standard “domain decomposition” using METIS

<http://www.hpc.lsu.edu/docs/guides.php?system=QB2>






- ADCIRC meshes can be enormous
- Run in parallel on super computers
- Actually meshes are decomposed (partitioned into many smaller meshes)
- Standard method involves “METIS”, a library of algorithms aimed at mesh decomposition



<https://ccht.ccee.ncsu.edu/figuregen-v-49/>



# METIS Partitioning



# Why Explore Alt. Domain Decomposition?

- ◆ ADCIRC is parallelized using a “message passing” technique (MPI)
- ◆ All touching subdomains must communicate over the network, network saturation is a known scaling limitation with ADCIRC
- ◆ METIS does not optimize for *communication complexity*, so there's interest in exploring others that do

# Apply DBSCAN to a FE Mesh?

## Pros

- ◆ ADCIRC meshes are most dense in areas of interest
- ◆ May result in net fewer “touching” subdomains
- ◆ Algorithm is “simple” and can be optimized in several ways

(spoiler: using OpenMP)

## Cons

- ◆ Also doesn't minimize communication complexity directly
- ◆ Computational complexity is  $O(n^2)$
- ◆ A constant ( $\epsilon$ ) that defines points that are “close”
- ◆ It can be non-deterministic

# ***Resulting Clusters***

- ◆ ... see images
- ◆ Results are in the right direction, but not as well defined ... even accounting for the coarse application of  $\varepsilon$
- ◆ HDBSCAN is a variation of DBSCAN that efficiently computes a solution using variable  $\varepsilon$ , and allows for the reconstruction of all the desired clusters

# ***Code Inspection Time ...***

(remember to open the code in the terminal)

# Performance && Correctness

- ◆ Massive improvement moving from pure Perl to **Inline::C**
- ◆ No scaling improvements (see **NYTProf**) when using more threads, which is okay *for now*
- ◆ htop shows definitely that the threads are all working (show demo)
- ◆ Correctness of the clustering is suspect, more work to be done there also



# An Examination of Code

## The Good

- ◆ `Inline::C`, `Alien::OpenMP`, and `OpenMP::Environment` are all used
- ◆ Convenient macros were identified and *used* (e.g., `_POMP_ENV_UPDATE_NUM_THREADS_`)
- ◆ An appropriate execution module is being clarified
- ◆ The algorithm did kind of work as implemented, but some results are more muddled than I wanted
- ◆ I learned a lot about using the Perl API from the `Inline::C::Cookbook` (thank you Ingy, et al.!!)

# An Examination of Code The Bad

- ◆ Expected scaling improvements were not achieved,
- ◆ e.g., I could not bare to move the outter loop in the Perl code into the C, which would have allowed the use of OpenMP's ability to “collapse” nested loops across threads when correctness was not violated (like in this case).
- ◆ Useful idioms were not able to be identified

# An Examination of Code

## The Ugly

- ◆ Perl data structures are wholly insufficient
- ◆ The required Perl API coding was way too low level
- ◆ The resulting code was neither Perlish nor at appropriately high level for most users of OpenMP

# Confronting Ugly Pt. 1

The Perl API is Too Low Level for OpenMP

- ◆ Because Perl data types are internally so extremely uniprocess (which is actually a very powerful feature), and perhaps because of my own extreme knowledge gaps; I optioned to simply dump the data into basic C data types before any threading
- ◆ The Perl API code I wrote was very scary and very ugly



# Confronting Ugly Pt. 2

## Perl Needs Thread-safe Datatypes

- ◆ I had to use the Perl API code not only because I could not safely modify an HV\* or AV\* across thread (which I verified as unsafe), but ...
- ◆ also because I was assured the even reading these structures by competing threads was not safe (I have not verified this, but I trust this is true)



# The Beautiful Pt. 1

## Higher Level of Perl API functions

- ◆ *For dumping 1d, 2d, 3d AVs of a uniform type (int, float) into C dimensionally equivalent arrays (OpenMP's bread & butter!)*
- ◆ *Easier iterating over single and double leveled Hvs (even if values were constrained by C type)*
- ◆ *Way easier inspection of Hvs (keys to C array, sorting, and less tedious ways to “**hv\_fetch**”, get the value given the key*
- ◆ *Would be nice in as part of the core Perl API, but can likely be distributed with **Alien::OpenMP***



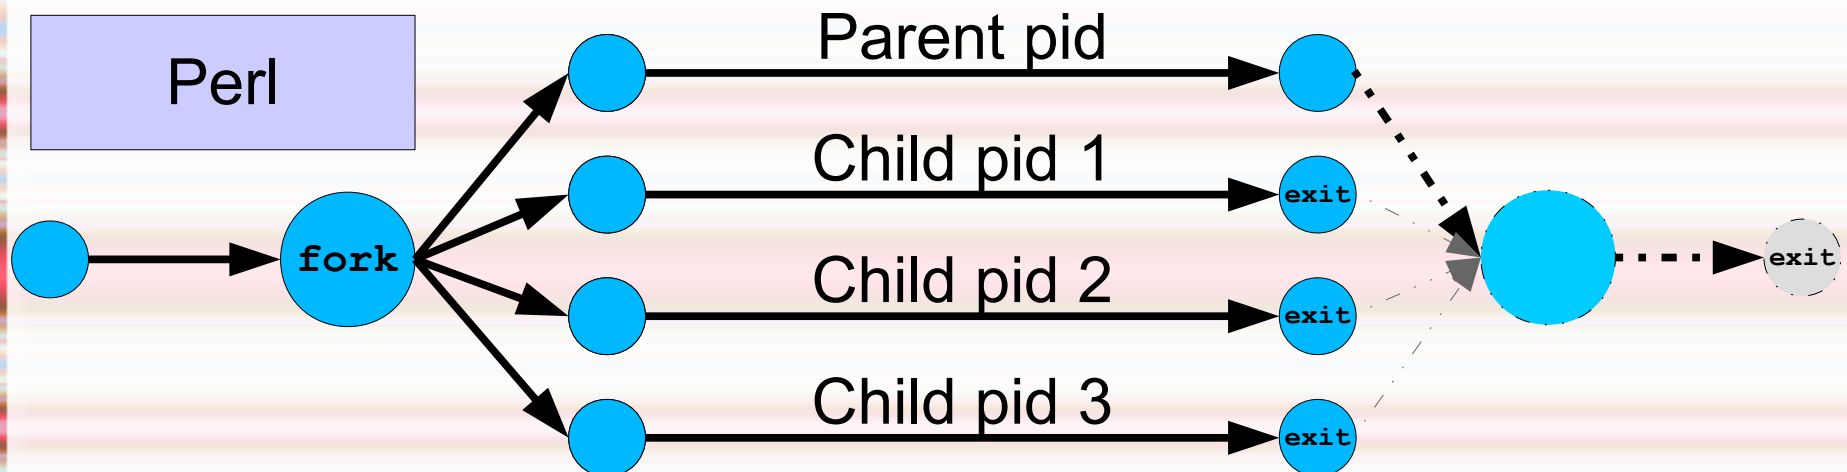
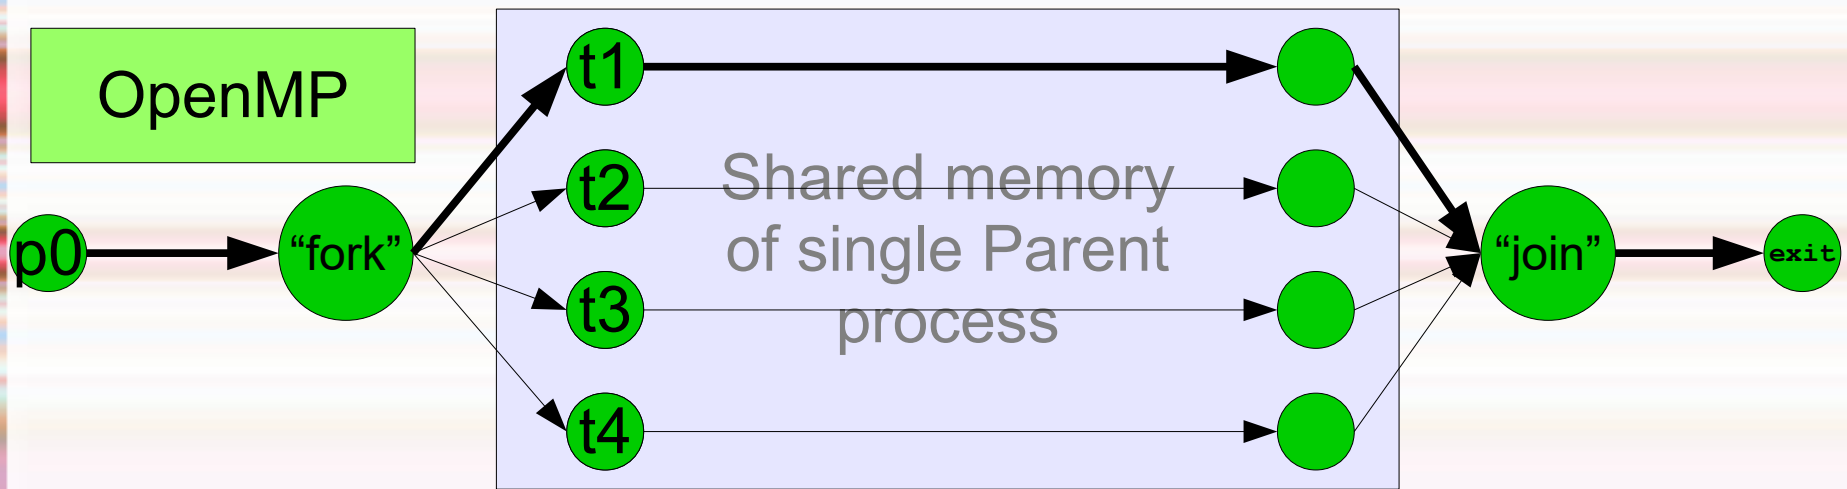
# The Beautiful Pt. 2

Simplified, thread-safe Data structures in Core

- ◆ Lockfree data structures seems to fit extensively well here (skip atomics, at least for now)
- ◆ AVs, HVs that can be modified, even in a limited manner by competing threads, even if they exist in Perl API-land only
- ◆ Easy conversion between the lockfree data structures and current types

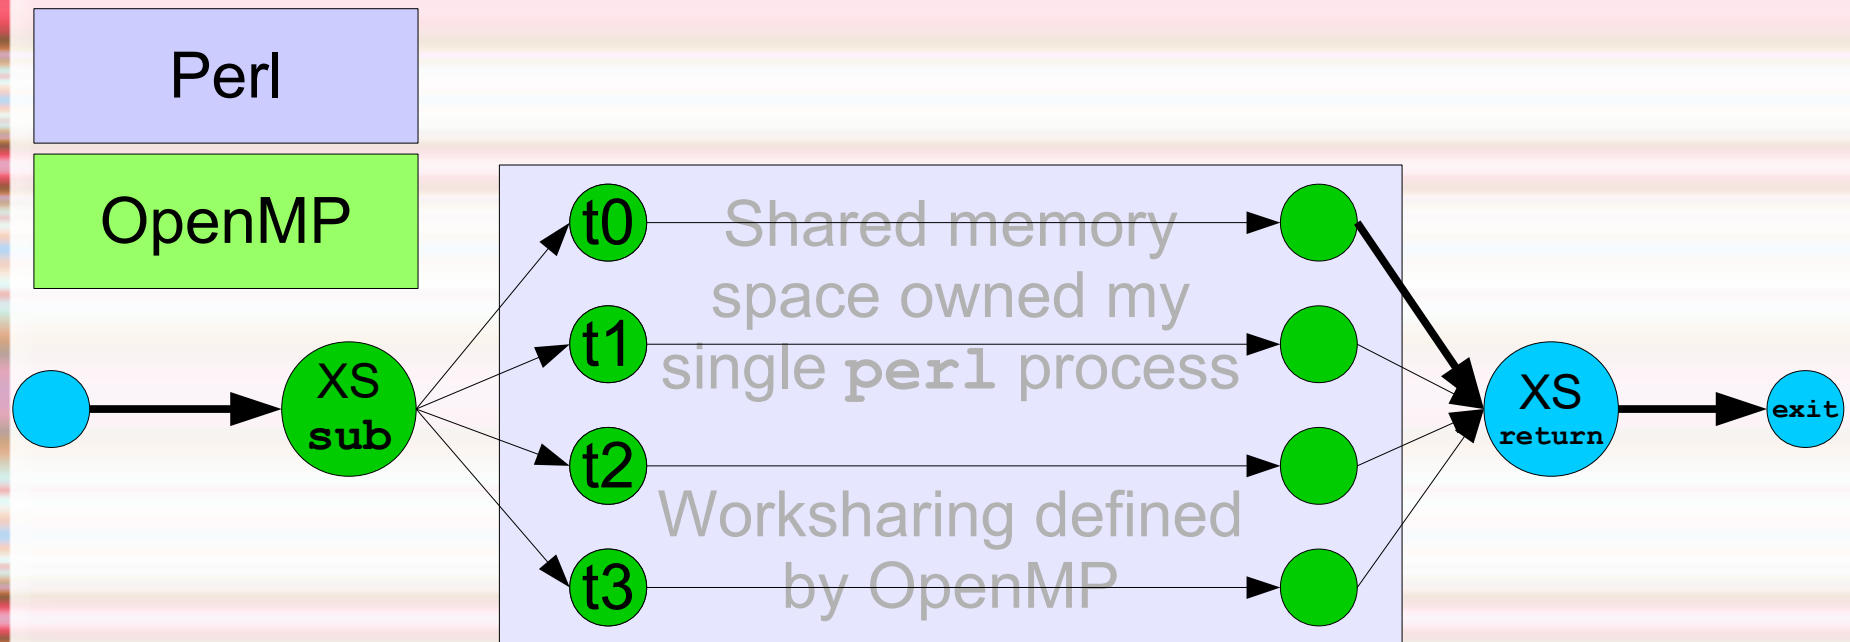
# The Beautiful Pt. 3a

Perl and OpenMP Execution Models are Complementary ...



# The Beautiful Pt. 3b

An Execution model Perl+OpenMP enables ...



# The Beautiful Pt. 3c

Perl+OpenMP w/ Perl Datastructure Support

- ◆ Largely avoids the need to get HV\* or AV\* data into normal C containers
- ◆ Avoids use of atomic or critical sections, which in OpenMP; pipelining critical sections with independent Perl structures helps but doesn't solve the issue (simultaneous reads are also dangerous)
- ◆ Lockfree datastructures are well know and extremely relevant
- ◆ Minimal support of new types would go a long way!

# Next Steps of Action

- ◆ Continuing working on the mesh decomposition tool using this method and others (e.g., k-part)
- ◆ Add a “helper” library **#include**, that contains useful MACROs (e.g., for **OMP\_NUM\_THREADS** envvar) and “missing” Perl API functionality to **Alien::OpenMP**
- ◆ I will continue to force myself to learn Perl API, one day I am sure I will love it
- ◆ If I can implement any or all of the above, I will surely try
- ◆ We need to do something; help and feedback is wanted!
- ◆ TIMTOWTDI in full effect (let's do this!)

Thanks, Eh!

